



POSTED BY

EXODUS INTEL VRT



POSTED ON

MARCH 20, 2019



POSTED UNDER

UNCATEGORIZED

## RECENT POSTS

Patch-gapping  
Google Chrome

Pwn2Own 2019:  
Microsoft Edge  
Sandbox Escape  
(CVE-2019-0938).  
Part 2

Pwn2Own 2019:  
Microsoft Edge  
Renderer  
Exploitation

## CVE-2019-5786: ANALYSIS & EXPLOITATION OF THE RECENTLY PATCHED CHROME VULNERABILITY

*This post provides detailed analysis and an exploit achieving remote code execution for the recently fixed Chrome vulnerability that was observed by Google to be exploited in the wild.*

Author: István Kurucsai

## Patch Analysis

The [release notes](#) from Google are short on information as usual:

[\$N/A][936448] High CVE-2019-5786: Use-after-free in FileReader. Reported by Clement Lecigne of Google's Threat Analysis Group on 2019-02-27

As described on [MDN](#), the “FileReader object lets web applications asynchronously read the contents of files (or raw data buffers) stored on the user's computer, using File or Blob objects to specify the file or data to read”. It can be used to read the contents of files selected in a file open dialog by the user or Blobs created by script code. An example usage is shown in below.

```
1 let reader = new FileReader();
2
3 reader.onloadend = function(evt) {
4   console.log(`contents as an ArrayBuffer: ${evt.target.result}`);
5 }
6
7 reader.onprogress = function(evt) {
8   console.log(`read ${evt.target.result.byteLength} bytes`);
9 }
10
11 let contents = "filecontents";
12 f = new File([contents], "a.txt");
13 reader.readAsArrayBuffer(f);
```

(CVE-2019-0940).  
Part 1

Windows Within  
Windows –  
Escaping The  
Chrome Sandbox  
With a Win32k  
NDay

A window of  
opportunity:  
exploiting a  
Chrome 1day  
vulnerability

### RECENT COMMENTS

Stagefright Patch  
Incomplete  
Leaving Android  
Devices Still  
Exposed | The  
Root Shell on

It is important to note that the File or Blob contents are read asynchronously and the user JS code is notified of the progress via callbacks. The *onprogress* event may be fired multiple times while the reading is in progress, giving access to the contents read so far. The *onloadend* event is triggered once the operation is completed, either in success or failure.

Searching for the issue number in the Chromium git logs quickly reveals the [patch](#) for the vulnerability, which alters a single function. The original, vulnerable version is shown below.

```
1 DOMArrayBuffer* FileReaderLoader::ArrayBufferResult() {
2   DCHECK_EQ(read_type_, kReadAsArrayBuffer);
3   if (array_buffer_result_)
4     return array_buffer_result_;
5
6   // If the loading is not started or an error occurs,
7   if (!raw_data_ || error_code_ != FileErrorCode::kOK)
8     return nullptr;
9
10  DOMArrayBuffer* result = DOMArrayBuffer::Create(raw_data_);
11  if (finished_loading_) {
12    array_buffer_result_ = result;
13    AdjustReportedMemoryUsageToV8(
14      -1 * static_cast<int64_t>(raw_data_>ByteLength());
15    raw_data_.reset();
16  }
17  return result;
18 }
```

This function gets called each time the result property is accessed in a callback after a *FileReader.readAsArrayBuffer* call in JavaScript.

Stagefright:

Mission

Accomplished?

---

Stagefright Patch

Incomplete

Leaving Android

Devices Still

Exposed |

Threatpost | The

first stop for

security news on

Stagefright:

Mission

Accomplished?

---

Tails live OS

affected by

critical zero-day

vulnerabilities on

Silver Bullets and

Fairy Tails

---

Are Tor and Tails

Safe? | SxiSpiGrl

on Silver Bullets

and Fairy Tails

---

While the object hierarchy around the C++ implementation of `ArrayBuffers` is relatively complicated, the important pieces are described below. Note that the C++ namespaces of the different classes are included so that distinguishing between objects implemented in Chromium (the `WTF` and `blink` namespaces) and `v8` (everything under the `v8` namespace) is easier.

- **`WTF::ArrayBuffer`**: the embedder-side (Chromium) implementation of the `ArrayBuffer` concept.  
**`WTF::ArrayBuffer`** objects are reference counted and contain the raw pointer to their underlying memory buffer, which is freed when the reference count of an `ArrayBuffer` reaches 0.
- **`blink::DOMArrayBufferBase`**: a garbage collected class containing a smart pointer to a **`WTF::ArrayBuffer`**.
- **`blink::DOMArrayBuffer`**: class inheriting from **`blink::DOMArrayBufferBase`**, describing an `ArrayBuffer` in Chromium. Represented in the JavaScript engine by a **`v8::internal::JSArrayBuffer`** instance.
- **`WTF::ArrayBufferBuilder`**: helper class to construct a **`WTF::ArrayBuffer`** incrementally. Holds a smart pointer to the `ArrayBuffer`.
- **`blink::FileReaderLoader`**: responsible for loading the File or Blob contents. Uses **`WTF::ArrayBufferBuilder`** to build

Scott Herbert  
(@Scott\_Herbert)  
on Silver Bullets  
and Fairy Tails

## ARCHIVES

---

September 2019

---

May 2019

---

April 2019

---

March 2019

---

January 2019

---

October 2018

---

September 2018

---

October 2017

---

July 2017



the `ArrayBuffer` as the data is read.

Comparing the code to the fixed version shown below, the most important difference is that if loading is not finished, the patched version creates new `ArrayBuffer` objects using the `ArrayBuffer::Create` function while the vulnerable version simply passes on a reference to the existing `ArrayBuffer` to the `DOMArrayBuffer::Create` function. `ToArrayBuffer` always returns the actual state of the `ArrayBuffer` being built but since the reading is asynchronous, it may return the same one under some circumstances.

```
1 DOMArrayBuffer* FileReaderLoader::ArrayBufferResult() {
2   DCHECK_EQ(read_type_, kReadAsArrayBuffer);
3   if (array_buffer_result_)
4     return array_buffer_result_;
5
6   // If the loading is not started or an error occurs,
7   if (!raw_data_ || error_code_ != FileErrorCode::kOK)
8     return nullptr;
9
10  if (!finished_loading_) {
11    return DOMArrayBuffer::Create(
12      ArrayBuffer::Create(raw_data_>Data(), raw_data_
13    )
14  }
15  array_buffer_result_ = DOMArrayBuffer::Create(raw_data_
16  AdjustReportedMemoryUsageToV8(-1 *
17                                static_cast<int64_t>(raw
18  raw_data_.reset();
19  return array_buffer_result_;
20 }
```

What are those circumstances? The `raw_data_` variable in the code is of the type `ArrayBufferBuilder`, which is used to construct

February 2017

January 2017

September 2016

August 2016

July 2016

June 2016

May 2016

February 2016

August 2015

April 2015

December 2014

August 2014

July 2014

December 2013

November 2013

the result `ArrayBuffer` from the incrementally read data by dynamically allocating larger and larger underlying `ArrayBuffers` as needed. The `ToArrayBuffer` method returns a smart pointer to this underlying `ArrayBuffer` if the contents read so far fully occupy the currently allocated buffer and creates a new one via slicing if the buffer is not fully used yet.

```
1 scoped_refptr<ArrayBuffer> ArrayBufferBuilder::ToArrayBu
2 // Fully used. Return m_buffer as-is.
3 if (buffer->ByteLength() == bytes_used_)
4     return buffer_;
5
6 return buffer_->Slice(0, bytes_used_);
7 }
```

One way to abuse the multiple references to the same `ArrayBuffer` is by detaching the `ArrayBuffer` through one and using the other, now dangling, reference. The javascript `postMessage()` method can be used to send messages to a JS Worker. It also has an additional parameter, *transfer*, which is an array of *Transferable* objects, the ownership of which are transferred to the Worker.

The transfer is done by the `blink::SerializedScriptValue::TransferArrayBufferContents` function, which iterates over the `DOMArrayBuffers` provided in the *transfer* parameter to `postMessage` and invokes the *Transfer* method of each, as shown below. *blink*:-

August 2013

---

January 2013

---

December 2012

---

November 2012

---

September 2012

---

August 2012

---

June 2012

## CATEGORIES

1Day

---

exploitation

---

internet explorer

---

NDay

---

News

---

*DOMArrayBuffer::Transfer* calls into *WTF::ArrayBuffer::Transfer*, which transfers the ownership of the underlying data buffer.

The vulnerability can be triggered by passing multiple *blink::DOMArrayBuffers* that reference the same underlying *ArrayBuffer* to *postMessage*. Transferring the first will take ownership of its buffer, then the transfer of the second will fail because its underlying *ArrayBuffer* has already been neutered. This causes

*blink::SerializedScriptValue::TransferArrayBufferContents* to enter an error path, freeing the already transferred *ArrayBuffer* but leaving a dangling reference to it in the second *blink::DOMArrayBuffer*, which can then be used to access the freed memory through JavaScript.

```
1 SerializedScriptValue::TransferArrayBufferContents(  
2 ...  
3 for (auto* it = array_buffers.begin(); it != array_buffers.end(); ++it) {  
4     DOMArrayBufferBase* array_buffer_base = *it;  
5     if (visited.Contains(array_buffer_base))  
6         continue;  
7     visited.insert(array_buffer_base);  
8  
9     wtf_size_t index = static_cast<wtf_size_t>(std::distance(array_buffers.begin(), it));  
10 ...  
11     DOMArrayBuffer* array_buffer = static_cast<DOMArrayBuffer*>(array_buffer_base);  
12  
13     if (!array_buffer->Transfer(isolate, contents.at(index), exception_state))  
14         exception_state.ThrowDOMException(DOMExceptionCode::INDEX_SIZE_ERR,  
15         "ArrayBuffer at index " +  
16         String::Number(index) +  
17         " could not be transferred.");  
18     return ArrayBufferContentsArray();  
19 }  
20 }
```

[Other](#)

[training](#)

[Uncategorized](#)

[Vulnerabilities](#)

META

[Log in](#)

[Entries RSS](#)

[Comments RSS](#)

[WordPress.org](#)

## Exploitation

The vulnerability can be turned into an arbitrary read/write primitive by reclaiming the memory region pointed to by the dangling pointer with JavaScript TypedArrays and corrupting their length and backing store pointers. This can then be further utilized to achieve arbitrary code execution in the renderer process.

## Memory Management in Chrome

There are several aspects of memory management in Chrome that affect the reliability of the vulnerability. Chrome uses **PartitionAlloc** to allocate the backing store of ArrayBuffers. This effectively separates ArrayBuffer backing stores from other kinds of allocations, making the vulnerability unexploitable if the region that is freed is below 2MiB in size because PartitionAlloc will never reuse those allocations for other kinds of data. If the backing store size is above 2MiB, it is placed in a directly mapped region. Once freed, other kinds of allocations can reuse such a region. However, successfully reclaiming the freed region is only possible on 32-bit platforms, as PartitionAlloc adds additional randomness to its allocations via VirtualAlloc and mmap address hinting on 64-bit platforms beside their ASLR slides.



On a 32-bit Windows 7 install, the address space of a fresh Chrome process is similar to the one shown below. Note that these addresses are not static and will differ by the ASLR slide of Windows. Bottom-up allocations start from the lower end of the address space, the last one is the reserved region starting at 36681000. Windows heaps, PartitionAlloc regions, garbage collected heaps of v8 and Chrome, thread stacks are all placed among these regions in a bottom-up fashion. The backing store of the vulnerable *ArrayBuffer* will also reside here. An important thing to note is that Chrome makes a 512MiB reserved allocation (from 4600000 on the listing below) early on. This is done because the address space on x86 Windows systems is tight and gets fragmented quickly, therefore Chrome makes an early reservation to be able to hand it out for large contiguous allocations, like *ArrayBuffers*, if needed. Once an *ArrayBuffer* allocation fails, Chrome frees this reserved region and tries again. The logic that handles this could complicate exploitation, so the exploit starts out by attempting a large (1GiB) *ArrayBuffer* allocation. This will cause Chrome to free the reserved region, then fail to allocate again, since the address space cannot have a gap of the requested size. While most OOM conditions kill the renderer process, *ArrayBuffer* allocation failures are recoverable from JavaScript via exception handling.

```

1  ...
2  45f5000 45f8000 3000 MEM_PRIVATE MEM_COMMIT PAGE_F
3  45f8000 4600000 8000 MEM_PRIVATE MEM_RESERVE
4  4600000 24600000 20000000 MEM_PRIVATE MEM_RESERVE
5  24600000 24601000 1000 MEM_PRIVATE MEM_COMMIT PAGE_
6  24601000 24602000 1000 MEM_PRIVATE MEM_RESERVE
7  ...
8  36681000 36690000 f000 MEM_PRIVATE MEM_RESERVE
9  36690000 65fc0000 2f930000 MEM_FREE PAGE_
10 65fc0000 65fc1000 1000 MEM_IMAGE MEM_COMMIT PAGE_
11 65fc1000 66085000 c4000 MEM_IMAGE MEM_COMMIT PAGE_
12 66085000 66086000 1000 MEM_IMAGE MEM_COMMIT PAGE_
13 ...

```

Another important factor is the non-deterministic nature of the multiple garbage collectors that are involved in the managed heaps of Chrome. This introduces noise in the address space that is hard to control from JavaScript. Since the *onprogress* events used to trigger the vulnerability are also fired a non-deterministic number of times, and each event causes an allocation, the final location of the vulnerable ArrayBuffer is uncontrollable without the ability to trigger garbage collections on demand from JavaScript. The exploit uses the code shown below to invoke garbage collection. This makes it possible to free the results of *onprogress* events continuously, which helps in avoiding out-of-memory kills of the renderer process and also forces the dangling pointer created upon triggering the vulnerability to point to the lower end of the address space, somewhere into the beginning of the original 512MiB reserved region.

```

1 function force_gc() {

```

```
2 // forces a garbage collection to avoid OOM kills and
3 try {
4   var failure = new WebAssembly.Memory({initial: 32767
5 } catch(e) {
6   // console.log(e.message);
7 }
8 }
```

## Exploitation steps

The exploit achieves code execution by the following steps:

- Allocate a large (128MiB) string that will be used as the source of the *Blob* passed to *FileReader*. This allocation will end up in the free region following the bottom-up allocations (from 36690000 in the address space listing above).
- Free the 512MiB reserved region via an oversized *ArrayBuffer* allocation, as discussed previously.
- Invoke *FileReader.readAsArrayBuffer*. A number of *onprogress* event will be triggered, the last couple of which can return references to the same underlying *ArrayBuffer* if the timing of the events is right. This step can be repeated indefinitely until successful without crashing the process.
- Free the backing store of the *ArrayBuffer* through one of the references. Going forward, another reference can be used to access the dangling pointer.

- Reclaim the freed region by spraying the heap with recognizable JavaScript objects, interspersed with TypedArrays.
- Look for the recognizable pattern through the dangling reference. This enables leaking the address of arbitrary objects by setting them as properties on the found object, then reading back the property value through the dangling pointer.
- Corrupt the backing store of a sprayed TypedArray and use it to achieve arbitrary read write access to the address space.
- Load a WebAssembly module. This maps a read-write-executable memory region of 64KiB into the address space.
- Traverse the JSFunction object hierarchy of an exported function from the WebAssembly module using the arbitrary read/write primitive to find the address of the read-write-executable region.
- Replace the code of the WebAssembly function with shellcode and execute it by invoking the function.

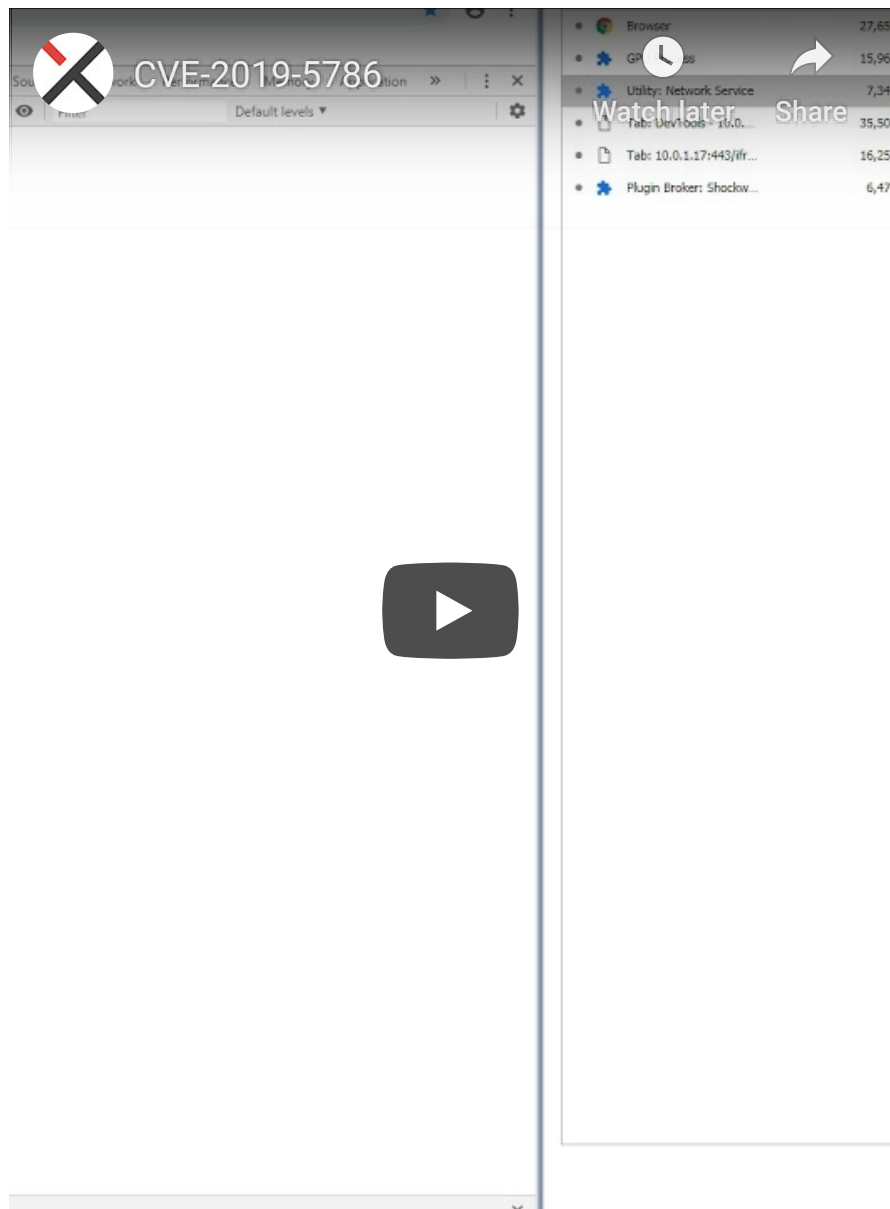
## Increasing reliability

A single run of the exploit (which uses the steps detailed above) yields a success rate of about 25%, but using a trick you can turn that into effectively 100% reliability. Abusing the site isolation



feature of Chrome enables brute-forcing, as described in [another post](#) on this blog by Ki Chan Ahn (look for the section titled “Making a Stealth Exploit by abusing Chrome’s Site Isolation”). A site corresponds to a (scheme:host) tuple, therefore hosting the brute forcing wrapper script on one site which loads the exploit repeatedly in an iframe from another host will cause new processes to be created for each exploit attempt. These iframes can be hidden from the user, resulting in a silent compromise. Using multiple sites to host the exploit code, the process can be parallelized (subject to memory and site-isolation process limits). The exploit developed uses a conservative timeout of 10 seconds for one iteration without parallelization and achieves code execution on average under half a minute.

The entire exploit code can be found on [our github](#) and it can be seen in action below.



## Detection

The exploit doesn't rely on any uncommon features or cause unusual behavior in the renderer process, which makes distinguishing between malicious and benign code difficult without false positive results.

## Mitigation

Disabling JavaScript execution via the Settings / Advanced settings / Privacy and security / Content settings menu provides effective mitigation against the vulnerability.

## Conclusion

It's interesting to see exploits in the wild still targeting older platforms like Windows 7 x86. The 32-bit address space is so crowded that additional randomization is disabled in PartitionAlloc and win32k lockdown is only available starting Windows 8. Therefore, the lack of mitigations on Windows 7 that are present in later versions of Windows make it a relatively soft target for exploitation.

Subscribers of our N-Day feed can leverage our in-depth analysis of critical vulnerabilities to defend themselves better, or use the provided exploits during internal penetration tests.

← Exploiting the Magellan bug on 64-bit Chrome Desktop

A window of opportunity: exploiting a Chrome 1day vulnerability →

---

Proudly powered by WordPress | Theme: Zoren by FabThemes.